

400,000 Ephemeral Containers

Testing entire ecosystems with Docker

Daniel Axtens

`linux.conf.au`, Hobart, 2017

- 125417 Ruby gems
- 369326 npm packages
- 95852 PyPI packages
- 9323 R packages

Test all the packages

- ① Simple
- ② Fast
- ③ Accurate

Test all the packages

- ① Simple
- ② Fast
- ③ Accurate

Does it install?

Let's do it

Live Demo!

Where we're going

1 Introduction and Motivation

2 Methodology

- Testing a package
- Managing a test
- Managing all the tests
- Going fast
- Display and post-processing

3 Review and results

- Results
- Interesting bugs

4 Future directions

- Other applications

5 Conclusion

- Final Thoughts

1 Introduction and Motivation

2 Methodology

- Testing a package
- Managing a test
- Managing all the tests
- Going fast
- Display and post-processing

3 Review and results

- Results
- Interesting bugs

4 Future directions

- Other applications

5 Conclusion

- Final Thoughts

Testing a package

- Tests shouldn't influence each other.
- Tests should be ephemeral.

- Isolation
- Low overhead
- Easy to script

- Virtual machines
- Chroots
- Containers

Live demo

Let's build a container...

Let's build a container...

```
FROM ppc64le/ubuntu:16.04
```

Let's build a container...

```
FROM ppc64le/ubuntu:16.04
```

```
ENV DEBIAN_FRONTEND=noninteractive
```

```
RUN apt-get update && \  
    apt-get upgrade -y && \  
    apt-get install -y build-essential curl libcurl4-openssl-dev \  
        libxml2-dev libpq-dev libmysqlclient-dev \  
        libpcre3-dev libxslt-dev libmagickcore-dev \  
        libmagickwand-dev libglib2.0-dev libv8-dev \  
        libpango1.0-dev default-jdk git cmake \  
        libmagic-dev libatk1.0-dev libsdl2-dev \  
        libsdl2-ttf-dev libpango1.0-dev libgsl0-dev \  
        libgl1-mesa-dev libfreeimage-dev \  
        libopenal-dev libsndfile-dev libgtk2.0-dev \  
        libgirepository1.0-dev libsasl2-dev \  
        libtag1-dev libavahi-compat-libdnssd-dev \  
        libmemcached-dev uuid-dev libsdl2-dev
```

Let's build a container...

```
RUN curl -sSL https://rvm.io/mpapis.asc | gpg --import - && \  
    curl -sSL https://get.rvm.io | bash -s stable --ruby=2.2.6 && \  
    rm -rf /var/cache/apt/* /usr/local/rvm/src && \  
    useradd -m -G rvm tester
```

Let's build a container...

```
RUN curl -sSL https://rvm.io/mpapis.asc | gpg --import - && \  
    curl -sSL https://get.rvm.io | bash -s stable --ruby=2.2.6 && \  
    rm -rf /var/cache/apt/* /usr/local/rvm/src && \  
    useradd -m -G rvm tester
```

```
USER tester
```


Let's build a container...

```
RUN curl -sSL https://rvm.io/mpapis.asc | gpg --import - && \  
    curl -sSL https://get.rvm.io | bash -s stable --ruby=2.2.6 && \  
    rm -rf /var/cache/apt/* /usr/local/rvm/src && \  
    useradd -m -G rvm tester
```

USER tester

```
RUN bash -c "source /usr/local/rvm/scripts/rvm && \  
    gem install --no-rdoc --no-ri nokogiri"
```

Let's build a container...

```
RUN curl -sSL https://rvm.io/mpapis.asc | gpg --import - && \  
    curl -sSL https://get.rvm.io | bash -s stable --ruby=2.2.6 && \  
    rm -rf /var/cache/apt/* /usr/local/rvm/src && \  
    useradd -m -G rvm tester
```

USER tester

```
RUN bash -c "source /usr/local/rvm/scripts/rvm && \  
    gem install --no-rdoc --no-ri nokogiri"
```

```
COPY gemrc /home/tester/.gemrc
```

```
ENV http_proxy=http://localhost:3128/ \  
    https_proxy=http://localhost:3128/
```

Let's build a container...

```
RUN curl -sSL https://rvm.io/mpapis.asc | gpg --import - && \  
    curl -sSL https://get.rvm.io | bash -s stable --ruby=2.2.6 && \  
    rm -rf /var/cache/apt/* /usr/local/rvm/src && \  
    useradd -m -G rvm tester
```

USER tester

```
RUN bash -c "source /usr/local/rvm/scripts/rvm && \  
    gem install --no-rdoc --no-ri nokogiri"
```

```
COPY gemrc /home/tester/.gemrc
```

```
ENV http_proxy=http://localhost:3128/ \  
    https_proxy=http://localhost:3128/
```

```
ENTRYPOINT ["/bin/bash", "-c"]
```

```
CMD ["source /usr/local/rvm/scripts/rvm && \  
    gem install --no-rdoc --no-ri $GEM"]
```

Live Demo!

1 Introduction and Motivation

2 Methodology

- Testing a package
- **Managing a test**
- Managing all the tests
- Going fast
- Display and post-processing

3 Review and results

- Results
- Interesting bugs

4 Future directions

- Other applications

5 Conclusion

- Final Thoughts

Managing tests...

```
#!/bin/bash
```

```
GEM=$1
```

```
TIMEOUT=240s
```

```
http_proxy=
```

```
echo "Testing $GEM"
```

```
cont=$(docker run -d --env GEM=$GEM --net=host \  
--tmpfs=/usr/local/rvm/gems/ruby-2.2.6/gems:rw,mode=777,gid=1000,exec \  
--tmpfs=/usr/local/rvm/gems/ruby-2.2.6/cache:rw,mode=777,gid=1000,exec \  
--tmpfs=/usr/local/rvm/gems/ruby-2.2.6/wrappers:rw,mode=777,gid=1000,exec \  
--tmpfs=/tmp/:rw,mode=777,exec \  
    rubytest)
```

Managing tests...

```
#!/bin/bash
```

```
GEM=$1
```

```
TIMEOUT=240s
```

```
http_proxy=
```

```
echo "Testing $GEM"
```

```
cont=$(docker run -d --env GEM=$GEM --net=host \  
--tmpfs=/usr/local/rvm/gems/ruby-2.2.6/gems:rw,mode=777,gid=1000,exec \  
--tmpfs=/usr/local/rvm/gems/ruby-2.2.6/cache:rw,mode=777,gid=1000,exec \  
--tmpfs=/usr/local/rvm/gems/ruby-2.2.6/wrappers:rw,mode=777,gid=1000,exec \  
--tmpfs=/tmp/:rw,mode=777,exec \  
    rubytest)
```

```
code=$(timeout $TIMEOUT docker wait "$cont")
```

Managing tests...

```
if [ -z "$code" ]; then
    echo "TIMEOUT on $GEM"
    docker kill $cont &> /dev/null
    docker logs $cont 2>&1 | curl -X PUT --data-binary @- \
        http://127.0.0.1/blaster/submit/ruby/${GEM}/T/
else
    if [ "$code" != "0" ]; then
        echo "FAILURE on $GEM"
        docker logs $cont 2>&1 | curl -X PUT --data-binary @- \
            http://127.0.0.1/blaster/submit/ruby/${GEM}/F/
    else
        echo "SUCCESS on $GEM"
        docker logs $cont 2>&1 | curl -X PUT --data-binary @- \
            http://127.0.0.1/blaster/submit/ruby/${GEM}/S/
    fi
fi
```


Managing tests...

```
if [ -z "$code" ]; then
    echo "TIMEOUT on $GEM"
    docker kill $cont &> /dev/null
    docker logs $cont 2>&1 | curl -X PUT --data-binary @- \
        http://127.0.0.1/blaster/submit/ruby/${GEM}/T/
else
    if [ "$code" != "0" ]; then
        echo "FAILURE on $GEM"
        docker logs $cont 2>&1 | curl -X PUT --data-binary @- \
            http://127.0.0.1/blaster/submit/ruby/${GEM}/F/
    else
        echo "SUCCESS on $GEM"
        docker logs $cont 2>&1 | curl -X PUT --data-binary @- \
            http://127.0.0.1/blaster/submit/ruby/${GEM}/S/
    fi
fi

docker rm $cont &> /dev/null
```

Let's test a package - with timeout

Live Demo!

1 Introduction and Motivation

2 Methodology

- Testing a package
- Managing a test
- **Managing all the tests**
- Going fast
- Display and post-processing

3 Review and results

- Results
- Interesting bugs

4 Future directions

- Other applications

5 Conclusion

- Final Thoughts

- `test.sh`

- test.sh

```
#!/bin/bash
```

```
parallel --joblog ruby.log --resume -j30 './testit.sh {}' ::: gems
```

- test.sh

```
#!/bin/bash
```

```
parallel --joblog ruby.log --resume -j30 './testit.sh {}' ::: gems
```

- gems

- test.sh

```
#!/bin/bash  
parallel --joblog ruby.log --resume -j30 './testit.sh {}' ::: gems
```

- gems

```
gem list --remote -q | awk '{print $1}' > gems
```

1 Introduction and Motivation

2 Methodology

- Testing a package
- Managing a test
- Managing all the tests
- **Going fast**
- Display and post-processing

3 Review and results

- Results
- Interesting bugs

4 Future directions

- Other applications

5 Conclusion

- Final Thoughts

Cache information for squid:

Hits as % of all requests:	5min: 65.8%, 60min: 66.7%
Hits as % of bytes sent:	5min: 64.5%, 60min: 72.0%

Other limiting factors

- Ruby — Disk IO:

```
--tmpfs=/usr/local/rvm/gems/ruby-2.2.6/gems:rw,mode=777,gid=1000,exec \  
--tmpfs=/usr/local/rvm/gems/ruby-2.2.6/cache:rw,mode=777,gid=1000,exec \  
--tmpfs=/usr/local/rvm/gems/ruby-2.2.6/wrappers:rw,mode=777,gid=1000,exec \  
--tmpfs=/tmp/:rw,mode=777,exec \
```

- NodeJS/npm

Other limiting factors

- Ruby — Disk IO:

```
--tmpfs=/usr/local/rvm/gems/ruby-2.2.6/gems:rw,mode=777,gid=1000,exec \  
--tmpfs=/usr/local/rvm/gems/ruby-2.2.6/cache:rw,mode=777,gid=1000,exec \  
--tmpfs=/usr/local/rvm/gems/ruby-2.2.6/wrappers:rw,mode=777,gid=1000,exec \  
--tmpfs=/tmp/:rw,mode=777,exec \
```

- NodeJS/npm

- Diminishing returns...

- 1 Introduction and Motivation
- 2 Methodology
 - Testing a package
 - Managing a test
 - Managing all the tests
 - Going fast
 - Display and post-processing
- 3 Review and results
 - Results
 - Interesting bugs
- 4 Future directions
 - Other applications
- 5 Conclusion
 - Final Thoughts

Excluding 'dumb stuff'

- Broken dependencies

```
ERROR: While executing gem ... (Gem::DependencyResolutionError)
    conflicting dependencies activesupport (= 3.0.0) \
        and activesupport (= 4.2.0)
```

- Out of date use of the C API

```
../fbuffer/fbuffer.h:175:47: error: macro "rb_str_new" requires 2
arguments, but only 1 given
    VALUE result = rb_str_new(FBUFFER_PAIR(fb));
```

- Other dumb errors/outright breakage

```
npm ERR! enoent ENOENT: no such file or directory, chmod \
    '/home/tester/node_modules/abn-validator/abn-validator.js'
npm ERR! enoent This is most likely not a problem with npm itself
npm ERR! enoent and is related to npm not being able to find a file.
npm ERR! enoent
```

1 Introduction and Motivation

2 Methodology

- Testing a package
- Managing a test
- Managing all the tests
- Going fast
- Display and post-processing

3 Review and results

- Results
- Interesting bugs

4 Future directions

- Other applications

5 Conclusion

- Final Thoughts


Summary of Results

Ruby > 96% success

NodeJS > 97% success

R \approx 98% success

Bugs: Packages



Docker Storage Drivers

- AUFS
 - Overheads:

```
Samples: 12M of event 'cycles', Event count (approx.): 4625267475995
overhead  shared object          symbol
 80.62%  [kernel]                [k] _raw_spin_lock
 1.35%   [kernel]                [k] find_busiest_group
```

- Random hangs
- Out-of-tree module
- Use Overlay

Bugs: Kernel (<https://lkml.org/lkml/2016/3/10/682>)

From Daniel Axtens <dja@axtens.net>

Subject ext4_file_open: Inconsistent encryption contexts (commit ff978b09f973) breaking Docker

Date Fri, 11 Mar 2016 11:44:54 +1100

Hi,

Trying to run a Docker container on a mainline kernel is failing intermittently, in interesting and exciting ways, such as:

```
$ docker run -it --rm --env PACKAGE=sinatra npmtest
operation not permitted
docker: Error response from daemon: Cannot start container \
4fc0120a6389f25241f84527a0d31854806f6fe4fd98d019f790cea0ae7e230b: \
[10] System error: operation not permitted.
```

...

1 Introduction and Motivation

2 Methodology

- Testing a package
- Managing a test
- Managing all the tests
- Going fast
- Display and post-processing

3 Review and results

- Results
- Interesting bugs


4 Future directions

- Other applications

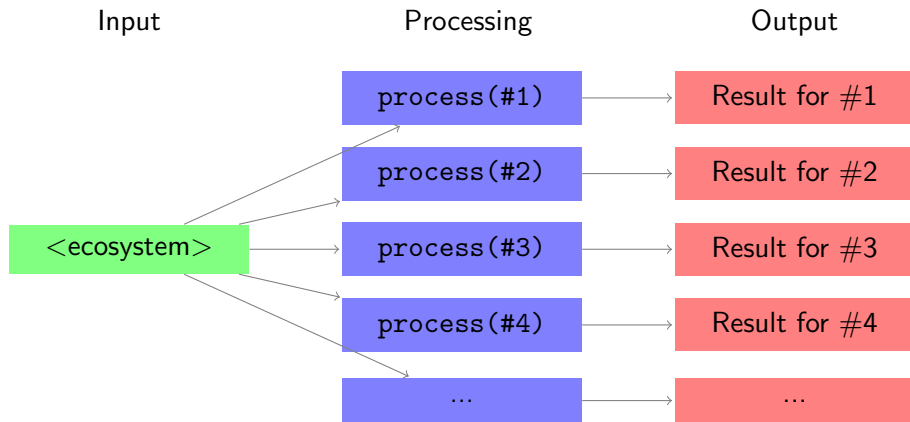
5 Conclusion

- Final Thoughts

Currently...



In general...



```
#include <stdlib.h>

int do_work(int value) {
    char * mem = malloc(1024);
    if (value < 0) {
        // we can't operate on negative values!
        return -1;
    }
    //do stuff with mem and value
    free(mem);
    return 42;
}
```

```
#include <stdlib.h>

int do_work(int value) {
    char * mem = malloc(1024);
    if (value < 0) {
        // we can't operate on negative values!
        return -1;
    }
    //do stuff with mem and value
    free(mem);
    return 42;
}
```

```
> cppcheck foo.c
Checking foo.c...
[foo.c:7]: (error) Memory leak: mem
```

s/gem install/cppcheck/

```
#!/bin/bash
cd /tmp
source /usr/local/rvm/scripts/rvm

gem fetch $GEM
gem unpack ${GEM}*.gem
cd ${GEM}*
if [ -d ext ]; then
    cppcheck --quiet --force ext/ || true
else
    echo "no source found for $GEM"
    exit 1
fi
```


s/gem install/cppcheck/

```
#!/bin/bash
cd /tmp
source /usr/local/rvm/scripts/rvm

gem fetch $GEM
gem unpack ${GEM}*.gem
cd ${GEM}*
if [ -d ext ]; then
    cppcheck --quiet --force ext/ || true
else
    echo "no source found for $GEM"
    exit 1
fi
```

The results are ... well, lets just do a live demo.

The offending code is:

```
static VALUE rb_gsl_blas_zgemm(int argc, VALUE *argv,
                               VALUE obj)
{
    gsl_matrix_complex *A = NULL, *B = NULL, *C = NULL;
    gsl_complex *pa = NULL, *pb = NULL;
    CBLAS_TRANSPOSE_t TransA, TransB;
    int flag = 0;
    (*pa).dat[0] = 1.0; (*pa).dat[1] = 0.0;
    (*pb).dat[0] = 0.0; (*pb).dat[1] = 0.0;
```

Recap

1 Introduction and Motivation

2 Methodology

- Testing a package
- Managing a test
- Managing all the tests
- Going fast
- Display and post-processing

3 Review and results

- Results
- Interesting bugs

4 Future directions

- Other applications

5 Conclusion

- Final Thoughts

Don't do it yourself

This project is a clumsy combination of other people's tools.

- GNU Parallel
- timeout
- squid
- Docker

The future of open source is...

Cool new stuff

Future of Open Source

The future of open source is...


~~Cool new stuff~~

The future of open source is...


Boring stuff at scale

- This work represents the view of the author and does not necessarily represent the view of IBM.
- IBM is a registered trademark of International Business Machines Corporation in the United States and/or other countries.
- Linux is a registered trademark of Linus Torvalds.
- Other company, product, and service names may be trademarks or service marks of others.

CI: Dependency Graph



CI: Reverse Dependency Graph



CI: Reverse Dependency Mapping

- Snap: []
- Crack: []
- Crackle: [Snap]
- Pop: [Snap, Crack]
- Fizz: [Snap]
- Buzz: [Fizz, Snap]